

Neural Network Framework



$$y = w_1x + b$$

CPU

Deep Learning Frameworks

- Scale ML code
- Compute Gradients!
- Standardize ML applications for sharing
- Interface with GPUs for parallel processing

Confusion Matrix

	PN	Yes	No
YN	TP	FN	
NO	FP	TN	

2×2 matrix

TensorFlow, Keras, and PyTorch



TensorFlow is an end-to-end open source platform for machine learning. It has a comprehensive, flexible ecosystem to build and deploy ML powered applications.



Keras is a high-level neural networks API, written in Python and capable of running on top of TensorFlow, CNTK, or Theano. It was developed with a focus on enabling fast experimentation.



PyTorch is an open source machine learning **framework** that accelerates the path from research prototyping to production deployment.

A Brief History of TensorFlow


TensorFlow is an end-to-end FOSS (free and open source software) library for dataflow, differentiable programming. TensorFlow is one of the most popular program frameworks for building machine learning applications.

- Google Brain built DistBelief in 2011 for internal usage.
- TensorFlow 1.0.0 was released on Feb 11, 2017
- TensorFlow 2.0 was released in Jan 2018.

What is TensorFlow?

TensorFlow allows for the specification and optimization of complex feed-forward models.

In particular, TensorFlow automatically differentiates a specified model.



A motivating example

The NumPy approach

```
net1 = W1 @ x + b1
h = s(net1)
net2 = W2 @ h + b2
output = s(net2)

# Manually compute derivative for W2
dL_doutput = (output - label)
dL_dW2 = dL_doutput @ ds(net) @ h
W2 -= learning_rate * dL_dW2
# Repeat for all other variables :)
```

The TensorFlow approach

```
net1 = W1 @ x + b1
h = tf.nn.sigmoid(net1)
net2 = W2 @ h + b2
output = tf.nn.sigmoid(net2)

# Let tensorflow do the heavy lifting
optimizer = tf.train.AdamOptimizer(learning_rate)
train = optimizer.minimize(I)

# Done :)
```

What is a Tensor in TensorFlow?

Rank - Unit of dimensionality described within the tensor

Shape - Number of rows and columns define together

- TensorFlow uses a tensor data structure to represent all data. A TensorFlow tensor has an n-dimensional array or list. A tensor has a static type, a rank, and a shape.

Name	Rank	Tensor
Scalar	0	[5]
Vector	1	[1 2 3]
Matrix	2	[[1 2 3 4], [5 6 7 8]]
Tensor	3	...

→ Row vector
Column vector

Type - Data type assigned to Tensor elements.

TensorFlow Data Types

Basic TensorFlow data types include:

- int[8|16|32|64], float[16|32|64], double
- bool
- string

with tf.cast(), the data types of variables could be converted.

TensorFlow Variables


TensorFlow variables can represent shared, persistent state manipulated by your program. Weights and biases are usually stored in variables.

```
import tensorflow as tf


W = tf.Variable(tf.random.normal([2,2], stddev=0.1),
                name = "W")
b = tf.Variable(tf.zeros(shape=(2)), name="b")
```

Keras is the official high-level API of TensorFlow

- tensorflow.keras (tf.keras) module
- Part of core TensorFlow since v1.4
- Full Keras API
- Better optimized for TF
- Better integration with TF-specific features




Machine Learning Workflow with tf.keras



What's special about Keras?

- A focus on user experience.
- Large adoption in the industry and research community.
- Multi-backend, multi-platform.
- Easy productization of models.



Industry Use



The Sequential API

```
import keras
from keras import layers

model = keras.Sequential()
model.add(layers.Dense(20, activation='relu', input_shape=(10,)))
model.add(layers.Dense(20, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.fit(x, y, epochs=10, batch_size=32)
```

The functional API

```
import keras
from keras import layers

inputs = keras.Input(shape=(10,))
x = layers.Dense(20, activation='relu')(x)
x = layers.Dense(20, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)

model = keras.Model(inputs, outputs)
model.fit(x, y, epochs=10, batch_size=32)
```

Model Subclassing

```
import keras
from keras import layers

class MyModel(keras.Model):

    def __init__(self):
        super(MyModel, self).__init__()
        self.dense1 = layers.Dense(20, activation='relu')
        self.dense2 = layers.Dense(20, activation='relu')
        self.dense3 = layers.Dense(10, activation='softmax')

    def call(self, inputs):
        x = self.dense1(x)
        x = self.dense2(x)
        return self.dense3(x)

model = MyModel()
model.fit(x, y, epochs=10, batch_size=32)
```

tf.keras Built-in Datasets

- tf.keras provides many popular reference datasets that could be used for demonstrating and testing deep neural network models. To name a few,
 - Boston Housing (regression)
 - CIFAR100 (classification of 100 image labels)
 - MNIST (classification of 10 digits)
 - Fashion-MNIST (classification of 10 fashion categories)
 - Reuters News (multiclass text classification)


Tensorflow

Create a tf.keras Model

- Layers are the fundamental building blocks of **tf.keras** models.
- The **Sequential** model is a linear stack of layers.
- A **Sequential** model can be created with a list of layer instances to the constructor or added with the **.add()** method.
- The input shape/dimension of the first layer need to be set.

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense,
Activation

model = Sequential([
    Dense(64, activation='relu', input_dim=20),
    Dense(10, activation='softmax')
])
```



Input Hidden Layers Output

Compile a tf.keras Model

The **compile** method of a Keras model configures the learning process before the model is trained. The following 3 arguments need to be set (the optimizer and loss function are required).

- An optimizer: **Adam**, **AdaGrad**, **SGD**, **RMSprop**, etc.
- A loss function: **mean_squared_error**, **mean_absolute_error**, **mean_squared_logarithmic_error**, **categorical_crossentropy**, **kullback_leibler_divergence**, etc.
- A list of measurement metrics: **accuracy**, **binary_accuracy**, **categorical_accuracy**, etc.